

Lumberyard Demo Setup Instructions

Author: Josh Patterson
jpatterson@floe.tv

Instructions

- Download VMWare
- Download a VM image
 - <http://www.thoughtpolice.co.uk/wmware>
 - Centos 5.5 is a great choice, 32bit
 - You may want to crank up memory to 2GB for the instance.
- Start your VM in VMWare
- Install Hadoop 0.20 on your VM in pseudo-distributed mode.
 - To make it easy, we'll use CDH3
 - <https://wiki.cloudera.com/display/DOC/Hadoop+%28CDH3%29+Quick+Start+Guide>
 - Easy way to do this:
 - `yum -y install hadoop-0.20-conf-pseudo`
 - You will need to upgrade Java (as talked about in the instructions)
 - Will also need to add Cloudera's repository
- Install HBase 0.89
 - To make it easy, we'll use Cloudera's repository
 - <https://wiki.cloudera.com/display/DOC/HBase+Installation>
 - Follow instructions closely
- Build jmotif from trunk at
 - <http://code.google.com/p/jmotif/>
 - A build of this is already included in the lib directory
 - We need this for the latest iSAX indexing commits
 - provides the bulk of the mathematics of the indexing scheme.
 - If you want to build jmotif yourself, you'll need at least ant 1.7 and you'll need to set the following env vars
 - `export JUNIT_HOME=<junit_home>`
 - `export FINDBUGS_HOME=<findbugs_home>`
 - `export WEKA_HOME=<weka_home>`
 - you'll need a later weka, preferably 3.7.2
 - To build the source:
 - `ant`
 - To build the jmotif jar (already one pre-built in the lib dir)
 - `ant -f jar.build.xml`
- Use Git to pull down Lumberyard from github

- git clone [git://github.com/jpatanooga/Lumberyard](https://github.com/jpatanooga/Lumberyard)
- Build
 - use ant to build the Lumberyard jar
 - ant
- Once we have both jmotif.lib.jar and Lumberyard.jar built, copy both jars into the VM onto the hbase path somewhere.
- To execute the Lumberyard shell, type:
 - **hbase tv.floe.lumberyard.hbase.isax.index.shell.Shell** <command>
 - To view the help
 - **hbase tv.floe.lumberyard.hbase.isax.index.shell.Shell**
 - To create an index
 - **hbase tv.floe.lumberyard.hbase.isax.index.shell.Shell**
CreateIndex "my_index"
 - -base_card 4
 - -dim_split 1
 - -base_word_len 8
 - -base_ts_sample_size 16
 - -split_threshold 100
 - Now let's index one of the two DNA files in the data subdir:
 - in the data subdirectory, we have the mitochondrial DNA of both the polar bear and the Hippo. These samples were downloaded from <http://www.cs.ucr.edu/~eamonn/iSAX/iSAX.html>
 - For a technical demo, let's search our index:
 - **hbase tv.floe.lumberyard.hbase.isax.index.shell.Shell**
SearchGenomeIndex "polar_bear" "TGGACTAAATGACTAA"
 - which should find one result
 - If we search for a slightly different pattern, we'll see its not found:
 - **hbase tv.floe.lumberyard.hbase.isax.index.shell.Shell**
SearchGenomeIndex "polar_bear" "TGGACTAAATGACTAG"
- The current demo
 - is built to demonstrate iSAX indexing and iSAX indexing persisted in HBase.
 - is meant to build indexes from a single thread
 - parallel index construction from Map Reduce is coming, requires some distributed lock coordination with something like Zookeeper
 - should not be counted on for more than prototyping purposes
 - is a proof of concept for scalable index construction and search
- Potential Use Cases
 - Genomic pattern matching
 - General timeseries search engine
 - Use variations with Map Reduce
 - Audio search
 - Image pattern search

References

- Based on the original paper by Shieh and Keogh:
 - <http://www.cs.ucr.edu/~eamonn/iSAX.pdf>
- iSAX homepage:
 - <http://www.cs.ucr.edu/~eamonn/iSAX/iSAX.html>